
MODULE 4

GETTING STARTED WITH R

Contents

4.1	Setting Up R and Helpers	23
4.2	Working With R Basics	23
4.3	Working With Data	25

4.1 Setting Up R and Helpers

Detailed methods for downloading, installing, and configuring R, RStudio, and `NCStats` on your personal computer are given on the [Resources page of the course website](#).

4.2 Working With R Basics

4.2.1 Saving Results

Results are not saved in R or RStudio. Rather, “scripts” of successful R commands are saved and, then, if the analysis needs to be re-done, the entire set of commands is opened in RStudio and run again. When writing a report, all tabular and graphical output should be copied from RStudio and pasted into your report document. This document will serve as your analysis report and can be modified to include answers to questions, references to the tables and graphs, etc.¹ All data that is not a simple vector (see Section 4.3.4) should be entered into R through text files (see Section 4.3.2).

R does allow one to save a “workspace”, though I urge you not to do that. Rather, save your “good” commands in a script and save your “good” results in a report document; do not save the workspace.

¹Specifics for how to format homework assignments is on the course syllabus

◇ Do NOT save the workspace in R.

4.2.2 Expressions and Assignments

Expressions in R are mathematical “equations” that are evaluated by R with a result seen immediately. An example of an expression in R is

```
> 5+log(7)-pi
[1] 3.804317
```

where `log()` and `pi` are built-in functions used to compute the natural log and find the value of π , respectively. Expressions in R are like using a calculator where the result is shown, but not saved for subsequent analyses. In addition, expressions in R follow the same order of operations and use of parentheses as expressions entered into your calculator.

◇ The results of expressions in R are temporary unless the result is assigned to an object.

Results from an expression are typically saved for further computations by assigning the results of the expression to an object with the assignment operator (i.e., `<-`). The general form for saving the result of an expression into an object is `object <- expression`. The result of the expression will not be seen unless the object name is subsequently typed into R (but see below). For example, the result of the previous expression is saved into an object called `x` and then viewed with

```
> x <- 5+log(7)-pi
> x
[1] 3.804317
```

The result of an expression can be both assigned and printed by surrounding the command in parentheses. For example, the following assigns the result of the expression to `y` and prints the result.²

```
> ( y <- 15*exp(2) )
[1] 110.8358
```

◇ The convention of surrounding commands in parentheses to both assign and print the results will be used extensively in this book to save space.


An object can be named whatever you want, with the exception that it cannot start with a number, contain a space, or be the name of a reserved word or function in R (e.g., `pi` or `log`). Furthermore, object names should be short and simple enough that you can remember what is contained in the object. It is also good practice to view the object immediately after making the assignment to make sure that it contains results that seem appropriate.


◇ In general, computational results should be assigned to an object.


²The spaces between the expression and the parentheses are only needed to increase legibility.


◇ Type the name of the object after making the assignment to confirm the results.


Review Exercises

4.1  Compute the value of $\frac{3}{7} + \frac{1}{2}$. [Answer](#)

4.2  Compute the value of $\pi * 3.7^2$. [Answer](#)

4.3  Assign the value of 3.7 to r . [Answer](#)

4.4  Compute the value of πr^2 using the value of r assigned in the previous problem. [Answer](#)

4.5  Assign the value 1.2 to r and then re-evaluate πr^2 . [Answer](#)

4.2.3 Functions and Arguments

R contains many “programs,” or functions, to perform particular tasks. A function is “called” by typing the function name followed by open and closed parentheses. Arguments, which the function will use to perform its task, are contained within the parentheses. The `log()` function, used in the previous section, is an example of a function. The name of the function is `log` and the argument, the number for which to compute the natural log, is contained within the parentheses following the function name. Many other functions will be described below and in subsequent modules.

△ **Function:** An R program that performs a particular task.

△ **Argument:** A “directive” that is provided to a function. Arguments are contained within parentheses that follow the function name.

◇ **Regular curved parentheses have two primary uses in R: (1) to control order of operations in expressions (as with a calculator) and (2) to contain the arguments sent to a function.**

4.3 Working With Data

4.3.1 Data Types

Data in R will be designated as an integer (whole numbers), numeric (non-integer numerical values), character (strings), factor (group membership), or logical (TRUE/FALSE). The type of data largely dictates the type of analysis that can be performed. Data types will be discussed in more detail as needed. Note, however,

that the **factor** data type is a special case of the character data type, where the specific items describe the group to which an individual belongs. This description allows for certain analyses in later modules.

△ **Factor:** A special type of variable that identifies the group to which an individual belongs.

4.3.2 Entering Data

For real data (i.e., several variables from many individuals) it is most efficient to enter data into a comma-separated values (CSV) file and then import that file into R. Creating a CSV file with Microsoft Excel is described below, though there are other ways to create CSV files (see [FAQs on class webpage](#)). This explanation assumes that you have a basic understanding of Excel (or other spreadsheet softwares).

◇ **Realistic datasets are most efficiently entered into a comma-separated values (CSV) file in preparation for importing into R.**

The spreadsheet should be organized with variables in columns and individuals in rows, with the exception that the first row should contain variable names. The example spreadsheet below shows the length (cm), weight (kg), and capture location data for a small sample of Black Bears.

	A	B	C
1	length.cm	weight.kg	loc
2	139	110	Bayfield
3	138	60	Bayfield
4	139	90	Bayfield
5	120.5	60	Bayfield
6	149	85	Bayfield
7	141	100	Ashland
8	141	95	Ashland
9	150	85	Douglas
10	166	155	Douglas
11	151.5	140	Douglas
12	129.5	105	Douglas
13	150	110	Douglas

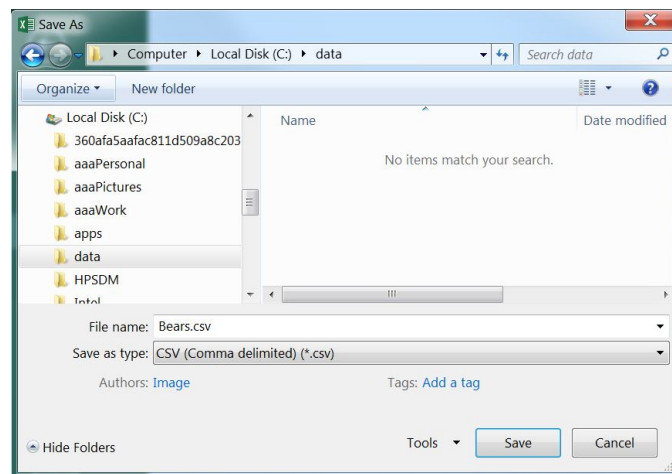
△ **data.frame:** A two-dimensional organization of variables (as columns, possibly of different data types) recorded on multiple individuals (as rows).

◇ **The columns of a data.frame correspond to variables and the rows of a data.frame correspond to individuals.**

Variable names must NOT contain spaces. For example, don't use *total length* or *length (cm)*. If you feel the need to have longer variable names, then separate the parts with a period (e.g., *length.cm*) or an underscore (e.g., *length_cm*). Furthermore, numerical measurements should NOT include units (e.g., don't use 7 cm). Finally, for categorical data, make sure that all categories are consistent (e.g., do not have a column that contains both *bayfield* and *Bayfield*).

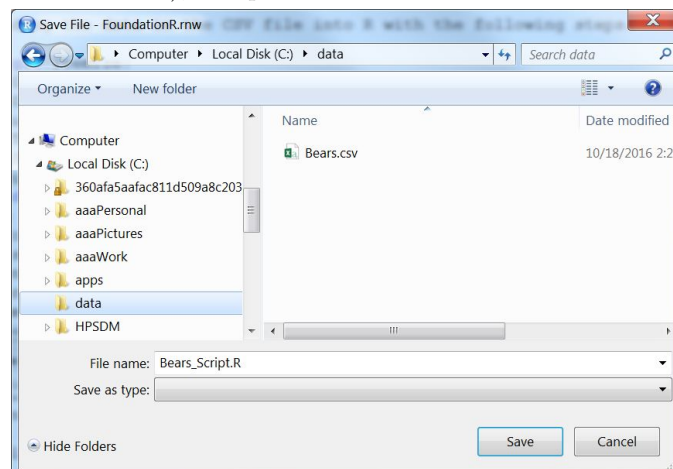
◊ **Variable names and data should not contain spaces. An "Error in scan" message usually indicates spaces in the variable names or data.**

The spreadsheet is saved as a CSV file by selecting the **File..Save As** menu item, which will produce the dialog box below. In this dialog box, change **Save as type** to **CSV (Comma delimited) (*.csv)** (you may have to scroll down), provide a file name (don't have any periods in the name besides for ".csv", which you should not have to type), select a location to save the file (don't forget this location!!), and press **Save**. Two "warning" dialog boxes may then appear – select **OK** for the first and **YES** for the second. You can now close the spreadsheet file (you may be asked to save changes – you should say **No**).



The following steps are used to load the data in the CSV file into RStudio.

- Open RStudio.
- Open a new script by selecting the **File, New File, R Script** menu items.
- Type **library(NCStats)** in the new script.
- Save this script by selecting the **File, Save** menu items. In the ensuing dialog box, navigate to the **exact same directory** where you saved the data, type a name for the file in the **File name:** box (**do not use a period in this name!!**), and press **Save**.



- Set the working directory (tell R where the file is) with the **Session, Set Working Directory ... To Source File Location** menu items in RStudio. RStudio will print an appropriate `setwd()` command to the console. Copy this command from the console to the second line in your script.³ For example, I stored the file created above in the `C:/data` directory, so that RStudio will create this `setwd("C:/data")`.
- The CSV file is read into R by including the name of the file (in quotes) in `read.csv()`. For example, `"Bears.csv"` is read into R and stored into an object called `bears` with `bears <- read.csv("Bears.csv")`.
- One should check the data in this object as described in Section 4.3.3 below

◊ Data stored in an external CSV file is read into R with `read.csv()`.

It is important that each row of the `data.frame` correspond to one individual. This is critically important when data are recorded for two different groups (e.g., for a two-sample t-test; see Module ??). For example, the following data are methyl mercury levels recorded in mussels from two locations labeled as “impacted” and “reference.”

```
impacted  0.011  0.054  0.056  0.095  0.051  0.077
reference  0.031  0.040  0.029  0.066  0.018  0.042  0.044
```

To follow the “one individual per row” rule, these data are entered in stacked format where the “reference” data are stacked underneath the “impacted” data and a column is used to indicate to which group the individuals belong. For example, the Excel file for data entry would look like the following

	A	B
1	loc	merc
2	impacted	0.011
3	impacted	0.054
4	impacted	0.056
5	impacted	0.095
6	impacted	0.051
7	impacted	0.077
8	reference	0.031
9	reference	0.04
10	reference	0.029
11	reference	0.066
12	reference	0.018
13	reference	0.042
14	reference	0.044

◊ Data files are constructed with data from only one individual in each row.

Alternative Forms of Getting Data

Some of the data files that you will use are provided on the [Data for MTH107](#) resource page of the class webpage. In these cases, the data should be downloaded from the webpage and saved in the same directory or folder as your analysis script. The downloaded file is then read into R in the same manner as described previously (i.e., set the working directory with `setwd()` and use `read.csv()`).

³Doing this will eliminate the need to manually select the menu options every time you want to run this script.

A few data files used in this book are supplied with R or the NCStats package. These files are loaded with `data()`. For example, the **iris** data file is loaded into R with

```
> data(iris)
```

4.3.3 Working With Data Frames

Viewing a Data Frame

Many users are disoriented in RStudio because they cannot “see” their data in the same way that they see it in a spreadsheet program. There are, however, several options for viewing your data. First, you can type the name of the data.frame object to see its entire contents.

```
> bears
  length.cm weight.kg      loc
1    139.0     110 Bayfield
2    138.0      60 Bayfield
3    139.0      90 Bayfield
4    120.5      60 Bayfield
5    149.0      85 Bayfield
6    141.0     100  Ashland
7    141.0      95  Ashland
8    150.0      85  Douglas
9    166.0     155  Douglas
10   151.5     140  Douglas
11   129.5     105  Douglas
12   150.0     110  Douglas
```

Typing the name is adequate for small data.frames, but not useful for large data.frames. The entire data.frame is opened in a separate window by double-clicking on the name of the data.frame in the **Environment** tab of RStudio. Alternatively, the first and last three rows are viewed by including the data.frame object in `headtail()`.

```
> headtail(bears)
  length.cm weight.kg      loc
1    139.0     110 Bayfield
2    138.0      60 Bayfield
3    139.0      90 Bayfield
10   151.5     140  Douglas
11   129.5     105  Douglas
12   150.0     110  Douglas
```

In addition to viewing the contents, it is useful to examine the structure of the data.frame as returned from `str()`. In this example, it is seen that three variables were recorded on 12 individuals. The first variables – *length.cm* and *weight.kg* – are numerical measurements made on the bears. The last variable – *loc* – is a factor variable that records the capture location for each bear.

```
> str(bears)
'data.frame': 12 obs. of 3 variables:
 $ length.cm: num  139 138 139 120 149 ...
 $ weight.kg: int  110 60 90 60 85 100 95 85 155 140 ...
 $ loc      : Factor w/ 3 levels "Ashland","Bayfield",...: 2 2 2 2 2 1 1 3 3 3 ...
```

The levels of the *loc* variable may be seen by including this variable (with the data.frame name) as the argument to `levels()`.

```
> levels(bears$loc)
[1] "Ashland" "Bayfield" "Douglas"
```

In the previous example, the `$` notation was used to identify a particular variable (i.e., *loc*) within a data.frame (*bears*). Think of variables as being nested inside data.frames and, thus, to access the variable you must first identify the data.frame in which it exists and then the name of the variable. The `$` simply separates the data.frame from the variable.

```
> bears$length.cm
[1] 139.0 138.0 139.0 120.5 149.0 141.0 141.0 150.0 166.0 151.5 129.5 150.0
> bears$loc
[1] Bayfield Bayfield Bayfield Bayfield Bayfield Ashland  Ashland  Douglas  Douglas
[10] Douglas  Douglas  Douglas
Levels: Ashland Bayfield Douglas
```

◇ A dollar sign is **ONLY** used in R to separate the name of a data.frame from the name of a variable within that data.frame.

Selecting Individuals

In some instances, it may be important to select or exclude an individual from a data.frame. Data.frames are two-dimensional objects that are indexed by a row and a column, in that order. Positions within a data.frame are selected within paired square brackets. For example, the item in the third row and second column of *bears* is selected below.

```
> bears[3,2]
[1] 90
```

◇ Identifying the position of an item in an object is the **ONLY** time that square brackets are used in R.

An entire row or column may be selected by omitting the other dimension. For example, one could select the entire second column with `bears[,2]`, but this is also the `weight.kg` variable and is better selected, as shown above, with `bears$weight.kg`. As a better example, the entire third row is selected below (note that the column designation was omitted).


```
> bears[3,]
  length.cm weight.kg    loc
3      139      90 Bayfield
```

Multiple rows are selected by combining row indices together with `c()` (more about `c()` in Section 4.3.4). For example, the third, fifth, and eighth rows are selected below (again, the column index is omitted).

```
> bears[c(3,5,8),]
  length.cm weight.kg    loc
3      139      90 Bayfield
5      149      85 Bayfield
8      150      85 Douglas
```

Finally, rows can be excluded by preceding the row indices with a negative sign.

```
> bears[-c(3,5,8,10,12),]
  length.cm weight.kg    loc
1      139.0      110 Bayfield
2      138.0       60 Bayfield
4      120.5       60 Bayfield
6      141.0      100 Ashland
7      141.0       95 Ashland
9      166.0      155 Douglas
11     129.5      105 Douglas
```

Filtering a data.frame

It is common to create a new data.frame that contains only some of the individuals from an existing data.frame. For example, a researcher may want a data.frame of only bears captured in Bayfield County or bears that weighed more than 100 kg. The process of creating the newer, smaller data.frame is called filtering (or subsetting) and is accomplished with `filterD()`. The `filterD()` function requires the original data.frame as the first argument and a condition statement as the second argument. The condition statement is used to either include or exclude individuals from the original data.frame. Condition statements consist of the name of a variable in the original data.frame, a comparison operator, and a comparison value (Table 4.1). The result from `filterD()` should be assigned to an object, which is then the name of the new data.frame.

◇ `filterD()` is used to create a new data.frame that consists of individuals selected by some criterion from an existing data.frame.

The following are examples of new data.frames created from *bears*. The name of the new data.frame (i.e., object to the left of the assignment operator) can be any valid object name. As demonstrated below, the new data.frame (or its structure) should be examined after each filtering to ensure that the data.frame actually contains the items that you desire.

- Only individuals from *Bayfield* county.

Table 4.1. Condition operators used in `filterD()` and their results. Note that *variable* generically represents a variable in the original data.frame and *value* is a generic value or level. Both *variable* and *value* would be replaced with specific items.

Condition Operator	Individuals Returned from Original Data Frame
<i>variable</i> == <i>value</i>	all individual that are equal to the given value
<i>variable</i> != <i>value</i>	all individuals that are NOT equal to the given value
<i>variable</i> > <i>value</i>	all individuals that are greater than the given value
<i>variable</i> >= <i>value</i>	all individuals that are greater than or equal to the given value
<i>variable</i> < <i>value</i>	all individuals that are less than the given value
<i>variable</i> <= <i>value</i>	all individuals that are less than or equal to the given value
<i>condition</i> , <i>condition</i>	all individuals that meet both conditions
<i>condition</i> <i>condition</i>	all individuals that meet one or both conditions ⁴

```
> bf <- filterD(bears,loc=="Bayfield")
> bf
  length.cm weight.kg      loc
1    139.0      110 Bayfield
2    138.0       60 Bayfield
3    139.0       90 Bayfield
4    120.5       60 Bayfield
5    149.0       85 Bayfield
```

- Individuals from both *Bayfield* and *Ashland* counties.

```
> bflash <- filterD(bears,loc %in% c("Bayfield","Ashland"))
> bflash
  length.cm weight.kg      loc
1    139.0      110 Bayfield
2    138.0       60 Bayfield
3    139.0       90 Bayfield
4    120.5       60 Bayfield
5    149.0       85 Bayfield
6    141.0      100  Ashland
7    141.0       95  Ashland
```

- Individuals with a weight greater than 100 kg.


```
> gt100 <- filterD(bears,weight.kg>100)
> gt100
  length.cm weight.kg      loc
1    139.0      110 Bayfield
2    166.0      155  Douglas
3    151.5      140  Douglas
4    129.5      105  Douglas
5    150.0      110  Douglas
```

- Individuals from *Douglas* County that weigh at least 150 kg.

```
> do150 <- filterD(bears,loc=="Douglas",weight.kg>=150)
> do150
  length.cm weight.kg    loc
1      166      155 Douglas
```

◇ View or “structure” the data.frame from filterD() to be sure that it contains data.

Review Exercises

4.6  Two students at Seattle Community College made biometric measurements on 25 Douglas fir (*Pseudotsuga menziesii*) trees in the lowlands of western Washington. The variables recorded in the [DougFirBiometrics.csv](#) file are a unique tree identifier (*tree*), the observer's name (*observer*; either “Ingrid” or “Dylan”), the circumference at breast height (meters; *circ*), the height to the eye of the observer (meters; *eyeht*), the horizontal distance between observer and tree (meters; *horizdist*), the angle between observer and top of tree (degrees; *angle*), and the estimated height of tree (meters; *height*) using right-angle trigonometry.

[Answer](#)

- Read this data file into an object called DF.
- Examine the structure of this data.frame.
- Show all measurements made on the third tree. [Do not do this manually; use R code.]
- Show all estimated tree heights.
- Show the estimated tree height for the fifth tree.
- Show all measurements for all trees measured by “Ingrid”. [HINT: use filtering.]
- Show all estimated tree heights for all trees measured by “Dylan”. [HINT: use filtering.]
- Show all measurements for tree heights less than 10 m. [HINT: use filtering.]
- Show all measurements for tree heights greater than 10 m and circumference less than 1 m. [HINT: use filtering.]

4.3.4 Vectors

Data.frames are the primary structure in which to store real data. However, much simpler situations that don't require a data.frame may arise. In R, items of the same data type (Section 4.3.1) are stored in a one-dimensional “list” called a *vector*. Vectors are usually displayed in one row (with many columns), but they may also be thought of as a single column (with many rows). Items are entered into a vector with `c()`, where the individual arguments are specific numbers, characters, or logical values.⁵ Items for a vector of characters must be contained within paired quotes.

⁵Note that `c` comes from the word “concatenate.”

```
> ( v <- c(1,2,5) )  
[1] 1 2 5  
> ( y <- c("Iowa","Minnesota","Wisconsin") )  
[1] "Iowa"      "Minnesota" "Wisconsin"
```






△ **Vector:** A one-dimensional list of items of the same data type. The primary information storage unit in R.

Single variables from a data.frame are vectors. Vectors that are not extracted from a data.frame will only be used in this course for very simple lists of items, usually as arguments in a function.

```
> bears$length.cm  
[1] 139.0 138.0 139.0 120.5 149.0 141.0 141.0 150.0 166.0 151.5 129.5 150.0
```

◇ The columns of a data.frame are accessed with the name of the data.frame, a dollar sign, and then the name of the variable – i.e., generically, data.frame\$varname.

Review Exercises

- 4.7  Create a vector called *h* that contains nine heights of people. [Answer](#)
- 4.8  Create a vector called *w* that contains nine weights of people. [Answer](#)
- 4.9  Create a vector called *hc* that contains nine hair colors of people. [Answer](#)
- 4.10  Create a vector called *m* that contains nine logical values (=TRUE if male). [Answer](#)
- 4.11  Using the vectors from the previous questions, [Answer](#)
- (a) ... create the largest possible data.frame (use data.frame()).
 - (b) ... identify the height of the third individual of this data.frame.
 - (c) ... identify the hair color for the sixth individual of this data.frame.